DEPARTMENT OF COMPUTER SCIENCE AT THE UNIVERSITY OF COPENHAGEN

Bachelors project

Autotuning Futhark

Simon Rotendahl (mpx651) & Carl Mathias Graae Larsen (pwh334)
{*simon, cala*}@di.ku.dk

Contents

June 2019

Abstract

1 Introduction

2 Background

GPGPU (General-purpose computing on graphics processing units) has been on the rise, due to limitation for CPU's, such as power limitation, and the fact that the computational potential of GPU's (Graphical processing unit) is vastly greater than CPU's [3] (a graph for theoretical computation, between the two, can be seen in B.1). There are technical/physical issue in reaching this potential, for example memory bandwidth, but there is also the conceptual issue of how to map parallelism to parallel hardware.

2.1 Parallelism

In the real world we have limitations to our current parallel hardware. We commonly execute parallel work on GPU's, due to its high number of threads. The limitation on GPU's, most essential to this report, is that a thread, may not spawn additional threads (this is not strictly true, but in practice it remains so). Therefore we need to work within that limitation. We say that since a GPU thread cannot spawn additional threads, it only supports *flat* parallelism. In the following sections we will look at approaches to work within this limitation.

2.2 Flat parallelism

Many problems does not have a flat structure, as our current platform effectively requires. Imagine a function (`drawline`) that, given two endpoints, will determine all the pixel that lie in a straight line in between. This can be done in parallel, but since only flat, and not nested parallelism is supported, we cannot call `drawline` within another parallel function (for example a `drawObject` function) [1].

Therefore we need to map nested parallelism to a flat structure. This problem has already been conceptually solved by *flattening*, and implemented in the language NESL [2]. The basic idea is to flatten the nested data that is to be worked on, and performing operations on that data in parallel.

When flattening a datastructure, we want to keep the information regarding the nest. In the original work, presenting flattening [1], this is done with a structure called *field* which has the form of (`nest information, value`), where a value can be an additional field, allowing for arbitrary depth. For a multidimensional array the field would be (`length, vaule/field`), below here is a concrete example.

`[[1,2,3],[4,5],[6,7,8,9]]` $\rightarrow$ `([3], ([3,2,4], [1,2,3,4,5,6,7,8,9]))`

This allows for a completely flat array, that can be mapped directly to the GPU. However the concept of flattening has practical issues. GPU's are not infinitely parallel, it has some capacity for parallelism,¡ and going beyond that will result in more overhead,

than computational gain. This problem is in no way addressed by flattening. Along with overhead, comes the issue that parallelism effectively destroys information about access patterns, making locality optimizations pretty much impossible, so we end up with greater overhead than gain, along with losing potential optimizations, when we always use the maximum parallelism available (fully flatten datastructures).

2.3 Futhark

The problem of mapping parallelism to flat hardware, and the difficulty of writing parallel code, in GPGPU languages such as CUDA and OpenCL, is what the programming language **Futhark** aims to solve. The creator of Futhark writes the purpose, of the language, nicely on the home page for the language *"Because it's nicer than writing CUDA or OpenCL by hand!"* [5]. On the same page, Futhark is described, more precisely, as *"a statically typed, data-parallel, and purely functional array language"*, but better than a description, is an example:

```
1 let dotprod [n] (xs: [n]f32) (ys: [n]f32): f32 =
2     reduce (+) 0f32 (map2 (*) xs ys)
3
4 let main [n][m][p] (xss: [n][m]f32) (yss: [m][p]f32): [n][p]f32 =
5     map (\xs -> map (dotprod xs) (transpose yss)) xss
```

Listing 1: Matrix-matrix multiplication in Futhark [6]

A Futhark program for matrix-matrix multiplication can be seen in listing 1, the syntax is similar to languages such as ML, and Haskell. It is a good example of how Futhark differs from CUDA or OpenCL (we would have liked to include an example of CUDA, but it was to long, so see A.1 for that). It allows the programmer to write efficient parallel code, without all the hardware specific knowledge regarding massively parallel systems.

The reason for Futhark being a functional language, is that functional languages lends themselves well to parallel problems, due to the functional paradigm not relying on evaluation order and side effects, as much as other many other paradigms. An example of this is `map`, which can process each element in parallel. However along with the problem of flat parallelism being required by the hardware, among other issues like to small of a stack and no function pointers, a functional language by itself is not a solution, and modifying an existing language such as Haskell is not optimal due to the size and expressiveness of it, being poorly suited for the restrictive nature of a GPU's.

In the following sections we will look at how Futhark deals with the problem of GPU's requiring flat parallelism. An example of a simple problem, that is not flat, is the `dotprod` function from listing 1. Here there is a `map2`[1] which is nested inside of a

---

[1]`map2` is a variant of `map` taking two arrays instead of a single element and an array. For example `map2 (+) [1,2,3] [4,5,6]` → `[5,7,9]`

reduce[2], both operations are parallel, thereby giving *nested parallelism.* In more general terms, nested parallelism is when there are parallel constructs nested within each other. Futhark supports such nested parallelism. More specifically, Futhark supports nested *data*-parallelism. Data-parallelism is when some function/operation, is applied to a dataset in parallel, such that one thread uses function x on a subset of data y, and another thread also uses function x, but on a different subset of y. Other languages/environments that support nested data-parallelism are the likes of; CUDA, Open MP, OpenCL, NESL etc. With NESL being the conceptual precursor to the other languages. As mentoned earlier, NESL solved the issue of nested data-parallelism by flattening data structures (both regular and irregular), resulting in maximum parallelism, which can be inefficient. Therefore we need to determine when it is useful to exploit more parallelism.

2.4 Moderate flattening

To determine the amount of parallelism to exploit, we need a closer loke at GPU's. In the world of GPU's we have *kernel*s, and a kernel is simply a GPU program. In CUDA specifically, it allows the programmer to define a function, that, when called, is executed N times in parallel by N different *CUDA threads* or *CUDA blocks*[3]. Such a function is called a kernel, and an example of a CUDA kernel is the matrix-matrix multiplication in appendix A.1.

A kernel can be thought of as the work, that is to be done, in a perfect *parallel nest.* A perfect parallel nest is a nest where all the work, is at the innermost level. Futhark transforms nested data-parallelism into structures, from which kernels can later

```
1 map (\xs -> let y = reduce (+) 0 xs
2         in map (+y) xs)
3    xss
4
```

```
1 let ys = map (\xs -> reduce (+) 0 xs)
       xss
2 in map (\xs y -> map (+y) xs) xss ys
3
```

Listing 2: Code fragment before distribution.

Listing 3: Code fragment after distribution.

Figure 1: Two version of a fragment of code, showing the distribution of SOACs (*second-order array combinators*) [4].

be extracted, and an example of this transformation can be seen in figure 1. We can see that the outer map, of listing 2, contains more parallelism in it, in terms of the parallel constructs reduce and map, so in the case where the outer map does not saturate the GPU's capacity, we would want to exploit the parallelism of the map body. With listing 2 this is not possible, since the outer map would be parallelised, and the inner sequentialized to one GPU kernel. Instead we can transform it to listing 3. Here the

---

[2]Reduce does not obviously appear to be a parallel construct, due to the elements being dependent on each other, however this is done by a *segmented reduce* [9], making it effectively so.

[3]For a further explanation of threads and blocks see [10, p. 5-9]

SOACs of the body (the parallel constructs `map` and `reduce`) in listing 2, are distributed out into their own `map` nests, giving two perfect `map` nests, and thereby translates to two GPU kernels, where the code of line 1 (eventual first kernel), is passed into the code of line 2 (eventual second kernel), allowing for further exploitation of the nested parallelism.

This transformation is dubbed *moderate flattening* [7], due to its conceptual resemblance to the flattening algorithm put forth by Blelloch and Sabot [1]. However it is moderate in its approach, due to it flattening the parallel construct until the parallelism saturates the hardware, after which it efficiently sequentialize the remaining parallelism. As we saw in Figure 1, Futhark flattens nested parallel construct by extracting kernels, based on rewrite/flattening rules. At the time of implementing moderate flattening, the algorithm used to apply these rules where based on heuristics about the structure of `map` nest contexts. These where decent approximations but due to the nature of GPU's (having different capacities for parallelism) there is no one size fits all.

It is important to note that Futhark does not currently support irregular data structures. This means that we cannot have a multidimensional array, where the elements have different shapes. The reason for this, is that to quantify an amount of parallelism, executed on a core efficiently, Futhark needs it to be regular. Irregular structures could be suspect to dynamic dependencies, and require more overhead because of the irregular memory access.

2.5   Incremental flattening

While moderate flattening was a good start it lacks the ability to dynamically flatten appropriately to the hardware capabilities, and the size of the data being worked on. This is where incremental flattening comes in. In short incremental flattening statically generates multiple different, semantically equivalent, piecewise code versions of the same program, based on the rules of moderate flattening (along with some additional ones) [8], and then dynamically chooses whether or not to further exploit parallelism or not.

To get an intuition for incremental flattening, lets go back to the matrix-matrix multiplication example in listing 1. Due to matrix-matrix multiplication containing a lot of `map`s (as many GPU programs do), we will briefly explain the code versions generated by the flattening rule, when a `map` containing additional parallelism, is encountered (CV is short for code version);

CV0 The body of the map will be executed sequentially. This means that there will be assigned one GPU thread to each element of the array being mapped over, each thread executing the map body sequentially.

CV1 The body of the map is partially executed in parallel. This means that there will be assigned one GPU group/block to each element of the array being mapped over. This is partial since it does not exploit the parallelism fully, due to, most likely, not having enough threads to fully parallelise the entire construct within the `map` body.

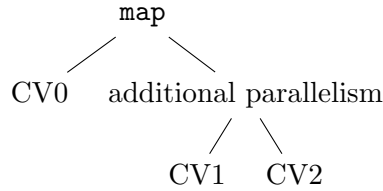CV2 Continues to flatten the `map` function, since we still have further parallel capacity to exploit.

```
                        map
                       /    \
                  CV0      additional parallelism
                                 /    \
                              CV1    CV2
```

Figure 2: A tree showing the structure of the three different code versions generate by flattening a `map` nest

In matrix-matrix multiplication, there are three levels of nested parallelism to exploit; the outer and inner `map` on line 5, and the `dotprod` function. Lets look at the different ways we can execute the function, there are at least 5 different ways to execute the code, generating 5 different code versions (denoted by a V) [8];

V0) The outer `map` is distributed out across the parallel construct, executing the body of the function sequentially. More specifically one GPU thread calculates one row in the resulting matrix each. This is CV0 for `map`.

V1) The outer `map` is executed in parallel, and the inner `map` of `main` is executed partially in parallel. More specifically one GPU group/block calculates one row in the resulting matrix. This is CV1 for `map`.

V2) The two outer maps of `main`, are extracted as kernels and the kernel of the outer `map` will invoke the inner `map` kernel saturating the GPU, and the body of the two maps (`dotprod`) is executed sequentially. More specifically each GPU thread is calculating one element in the resulting matrix. This is CV0 for `map`.

V3) The two outer `maps`, does not saturate the GPU, but the `dotprod` function would oversaturate it, therefore, `dotprod` is partially executed in parallel. More specifically each element of the result matrix is executed by a GPU group/block. This is CV1

V4) If the entire parallel nest does not oversaturate the GPU, the entire nest is executed in parallel. This is CV2

We have to choose between these different, but semantically equivalent, ways of execution. The optimal choice will depend on the hardware, and the data being worked on. To visualize the choices, we will look at them as a tree. In Figure 3 there are four choices, we have to consider, and these choices are dependent on each other

```
                    Choice0
                   /      \
                 V0      Choice1
                        /      \
                      V1      Choice2
                              /      \
                            V2      Choice3
                                    /      \
                                  V3       V4
```
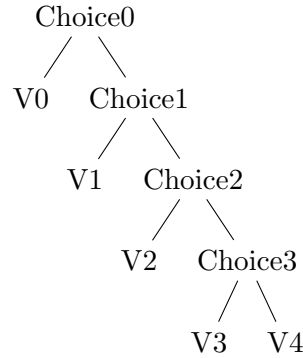
Figure 3: The structure of choices found in the futhark program for matrix-matrix multiplication (see listing 1). (V) represents the resulting code version of a choice

With these dependent choices, the idea of incremental flattening becomes increasingly clear. We step through the code, and decided whether we should flatten the parallelism available (for example a `map` nest), or whether we should sequentialize it, and exploit locality-of-reference optimizations. Due to the dependency of the choices, we end up iterating through the parallel nest, and flattening until we reach the full capacity of the hardware, hence the name *incremental flattening*. The choices will not always be in form of a unbalanced tree (although they are for the most part), they can also be balanced, and/or form a forest, we will give an example of a are more complicated tree later.

2.6   Structure of the code versions

To further understand how we will choose the correct code version, we have filled out the matrix-matrix multiplication tree, with more detail, which can be seen in Figure 4. In the tree, at the parent nodes, we see the statically generated predicates (the choice), hence forth known as *threshold comparisons*, or just *comparisons*, these guard the different code versions. The comparisons, have a *threshold parameter* on the left t_i, which is a symbolic representation of the parallel capacity of the hardware. On the right of the comparison, we have a value, based upon the dataset given, and as seen in the figure, these will reflect the shape and size of the data.
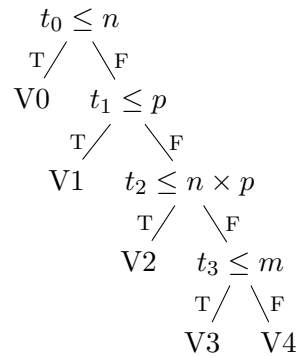
$$
\begin{array}{c}
t_0 \leq n \\
{\scriptstyle T}\diagup \quad \diagdown {\scriptstyle F} \\
V0 \qquad t_1 \leq p \\
\qquad {\scriptstyle T}\diagup \quad \diagdown {\scriptstyle F} \\
\qquad V1 \qquad t_2 \leq n \times p \\
\qquad\qquad {\scriptstyle T}\diagup \quad \diagdown {\scriptstyle F} \\
\qquad\qquad V2 \qquad t_3 \leq m \\
\qquad\qquad\qquad {\scriptstyle T}\diagup \quad \diagdown {\scriptstyle F} \\
\qquad\qquad\qquad V3 \qquad V4
\end{array}
$$

Figure 4: The tree generated by matrix-matrix multiplication. (t_i) is a threshold, $(n, p,$ and $m)$ are the sizes of the matricies, and (T) and (F) are indicative of the path we take, based on the threshold comparison.

To inspect the structure of these predicates and thresholds parameters further, lets look at a more complex Futhark program, from Futhark bench, called *LocVolCalib*;
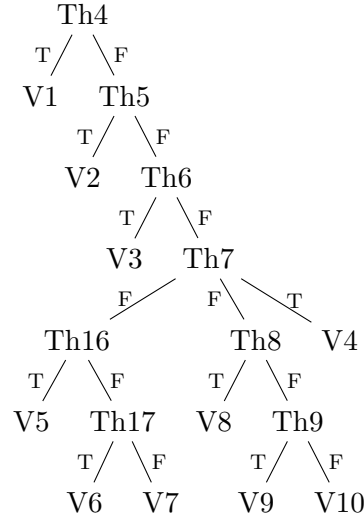

Figure 5: The dependencies between thresholds, of the test program `LocVolCalib.fut`. (Th) is a threshold comparison, (V) is a code version, and (T) and (F) are indicative of the path we take, based on the threshold comparison

To tune a program we need to examine each execution path through the tree. It is important not to get an end node (code version) confused with a version of the program. Two example of paths, and corresponding code versions, through the tree in Figure 5, that shows this, could be;

- $\{$(Th4, False), (Th5, False), (Th6, False), (Th7, True)$\} \rightarrow$ V4

- $\{$(Th4, False), (Th5, False), (Th6, False), (Th7, False), (Th8, False), (Th9, True), (Th16, False), (Th17, True)$\} \rightarrow$ (V6, V9)

The first path is simple, the code represented by `T4, T5, T6` is executed in parallel, where everything after it, is executed sequentially. The second path is more interesting, `T7` has two child nodes, that are reached with a false comparison. Here it is clear that two end nodes are reached, namely (`V6, V9`), and these two code versions are then combined into one program. This is also important to note, because we could have a forest, instead of a single tree, and this would leave multiple code versions, that are to be combined.

3   Search space and portability

There are two main reasons for the use of autotuning (compared to manual tuning), the search space of different threshold settings, and the portability of the code.

The search space   for thresholds is quite large. The threshold parameters has a value of 2^{15} as a default, lets us then consider matrix-matrix multiplication again (Figure 4). Here we have 4 thresholds, which would leave $\left(2^{15}\right)^4 \approx 1.153 \times 10^{18}$ different threshold parameter settings (and this is assuming that 2^{15} is an appropriate maximum value for all hardware, which is not the case). However we see that a setting can only end in five different executions. This is due to a data value of 128 would result in the same in all the settings from 2^{15} down 128. So to manually tune the search space is incredible tedious and repetitive.

Portability   of code is important, especially for GPU's due to the low level link between the code and the architecture (GPU threads, blocks/groups etc.). When tuning we tune to specific hardware, and representative datasets, so as soon as the code is to be executed on a different system (with a potentially different size of input data), the tuning performed would be useless (and in the worst case even detrimental to performance).

To illustrate the problem of repetitiveness through the search space we see Figure 4, with the thresholds $t = [10, 20, 30, 40]$, and the dataset values $dVal = [20, 30, 40, 50]$, so the predicate of $t_0 \leq dVal_0$ would result in taking the `true` branch, however this is also the case for $t_0 = 11$, making these two different settings, having the same dynamic behavior.

References

[1] Guy E. Blelloch and Gary W. Sabot. "Compiling Collection-oriented Languages Onto Massively Parallel Computers". In: *J. Parallel Distrib. Comput.* 8.2 (Feb. 1990), pp. 119–134. ISSN: 0743-7315. DOI: 10.1016/0743-7315(90)90087-6. URL: http://dx.doi.org/10.1016/0743-7315(90)90087-6.

[2] Guy E. Blelloch et al. "Implementation of a Portable Nested Data-parallel Language". In: *SIGPLAN Not.* 28.7 (July 1993), pp. 102–111. ISSN: 0362-1340. DOI: 10.1145/173284.155343. URL: http://doi.acm.org/10.1145/173284.155343.

[3] Michael Galloy. *CPU vs GPU performance.* June 2019. URL: https://michaelgalloy.com/2013/06/11/cpu-vs-gpu-performance.html. The graph used is an updated version from the article, and can be found here https://github.com/mgalloy/cpu-vs-gpu.

[4] Troels Henriksen. *Streaming Combinators and Extracting Flat Parallelism.* June 2019. URL: https://futhark-lang.org/blog/2017-06-25-futhark-at-pldi.html.

[5] Troels Henriksen. *Why Futhark?* May 2019. URL: https://futhark-lang.org/.

[6] Troels Henriksen et al. *Experimental infrastructure for the paper "Incremental Flattening for Nested Data Parallelism" at PPOPP'19.* May 2019. URL: https://github.com/diku-dk/futhark-ppopp19.

[7] Troels Henriksen et al. "Futhark: Purely Functional GPU-programming with Nested Parallelism and In-place Array Updates". In: *SIGPLAN Not.* 52.6 (June 2017), pp. 556–571. ISSN: 0362-1340. DOI: 10.1145/3140587.3062354. URL: http://doi.acm.org/10.1145/3140587.3062354.

[8] Troels Henriksen et al. "Incremental Flattening for Nested Data Parallelism". In: *Proceedings of the 24th Symposium on Principles and Practice of Parallel Programming.* PPoPP '19. Washington, District of Columbia: ACM, 2019, pp. 53–67. ISBN: 978-1-4503-6225-2. DOI: 10.1145/3293883.3295707. URL: http://doi.acm.org/10.1145/3293883.3295707.

[9] Rasmus Wriedt Larsen and Troels Henriksen. "Strategies for Regular Segmented Reductions on GPU". In: *Proceedings of the 6th ACM SIGPLAN International Workshop on Functional High-Performance Computing.* FHPC 2017. Oxford, UK: ACM, 2017, pp. 42–52. ISBN: 978-1-4503-5181-2. DOI: 10.1145/3122948.3122952. URL: http://doi.acm.org/10.1145/3122948.3122952.

[10] nVidia. "NVIDIA CUDA Compute Unified Device Architecture, Programming Guide". In: (2008). URL: http://developer.download.nvidia.com/compute/cuda/2_0/docs/NVIDIA_CUDA_Programming_Guide_2.0.pdf.

A   Code examples

A.1   Matrix Multiplication - CUDA for GPUs [10, p. 71-73]

```
__global__ void Muld(float* A, float* B, int wA, int wB, float* C)
{
    int bx = blockIdx.x;
    int by = blockIdx.y;
    int tx = threadIdx.x;
    int ty = threadIdx.y;

    int aBegin = wA * BLOCK_SIZE * by;
    int aEnd   = aBegin + wA - 1;
    int aStep  = BLOCK_SIZE;
    int bBegin = BLOCK_SIZE * bx;
    int bStep  = BLOCK_SIZE * wB;

    float Csub = 0;
    for (int a = aBegin, b = bBegin;
        a <= aEnd;
        a += aStep,
        b += bStep) {

        __shared__ float As[BLOCK_SIZE][BLOCK_SIZE];
        __shared__ float Bs[BLOCK_SIZE][BLOCK_SIZE];

        As[ty][tx] = A[a + wA * ty + tx];
        Bs[ty][tx] = B[b + wB * ty + tx];

        __syncthreads();

        for (int k = 0; k < BLOCK_SIZE; ++k)
            Csub += As[ty][k] * Bs[k][tx];
    __syncthreads();
    }

    int c = wB * BLOCK_SIZE * by + BLOCK_SIZE * bx;
    C[c + wB * ty + tx] = Csub;
}
```

B   Graphs

B.1   Theoretical performance for GPU's and CPU's [3]